

Assignment -4
SMS SPAM CLASSIFICATION

Assignment Date	21 October 2022
Student Name	Komala.P
Team ID	IBM-29269-1662616381
Maximum Marks	2 Marks

Import the Dataset

```
from google.colab import files  
uploaded = files.upload()
```

Saving spam.csv to spam.csv

Import required libraries

```
import csv  
import tensorflow as tf  
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
from tensorflow.keras.preprocessing.text import Tokenizer  
from tensorflow.keras.preprocessing.sequence import pad_sequences  
import nltk  
nltk.download('stopwords')  
from nltk.corpus import stopwords  
STOPWORDS = set(stopwords.words('english'))
```

[nltk_data] Downloading package stopwords to /root/nltk_data...

[nltk_data] Unzipping corpora/stopwords.zip.

[nltk_data] Unzipping corpora/stopwords.zip.

Import dataset

```
import io  
dataset = pd.read_csv(io.BytesIO(uploaded['spam.csv']), encoding = "ISO-8859-1")
```

dataset

```
      v1                                v2 Unnamed: 2 \
0   ham  Go until jurong point, crazy.. Available only ...   NaN
1   ham                Ok lar... Joking wif u oni...   NaN
2   spam  Free entry in 2 a wkly comp to win FA Cup fina...   NaN
3   ham  U dun say so early hor... U c already then say...   NaN
4   ham  Nah I don't think he goes to usf, he lives aro...   NaN
...   ...                                ...   ...
5567 spam  This is the 2nd time we have tried 2 contact u...   NaN
5568 ham                Will Ì_ b going to esplanade fr home?   NaN
5569 ham  Pity, * was in mood for that. So...any other s...   NaN
5570 ham  The guy did some bitching but I acted like i'd...   NaN
5571 ham                Rofl. Its true to its name   NaN
```

```
      Unnamed: 3 Unnamed: 4
0      NaN      NaN
1      NaN      NaN
2      NaN      NaN
3      NaN      NaN
4      NaN      NaN
...      ...      ...
5567     NaN     NaN
5568     NaN     NaN
5569     NaN     NaN
5570     NaN     NaN
5571     NaN     NaN
```

[5572 rows x 5 columns]

```
vocab_size =5000
embedding_dim =64
max_length =200
trunc_type ='post'
padding_type ='post'
oov_tok ="
training_portion =.8
```

Read the dataset and do pre-processing.

To remove the stop words.

```
articles = []
labels = []

with open("spam.csv", 'r', encoding="ISO-8859-1") as dataset:
    reader = csv.reader(dataset, delimiter=',')
    next(reader)
    for row in reader:
        labels.append(row[0])
        article = row[1]
        for word in STOPWORDS:
            token = ' ' + word + ' '
            article = article.replace(token, '')
            article = article.replace(' ', '')
        articles.append(article)
print(len(labels))
print(len(articles))

5572
5572
```

Train the model

```
train_size = int(len(articles) * training_portion)
train_articles = articles[0: train_size]
train_labels = labels[0: train_size]
validation_articles = articles[train_size:]
validation_labels = labels[train_size:]
print(train_size)
print(len(train_articles))
print(len(train_labels))
print(len(validation_articles))
print(len(validation_labels))

4457
4457
4457
1115
1115
```

```

tokenizer = Tokenizer(num_words = vocab_size, oov_token=oov_tok)
tokenizer.fit_on_texts(train_articles)
word_index = tokenizer.word_index
dict(list(word_index.items())[0:10])

{' ': 1,
 'i': 2,
 'u': 3,
 'call': 4,
 'you': 5,
 '2': 6,
 'get': 7,
 "i'm": 8,
 'ur': 9,
 'now': 10}

```

Training data to Sequences

```

train_sequences = tokenizer.texts_to_sequences(train_articles)
print(train_sequences[10])

[8, 190, 37, 201, 30, 260, 293, 991, 222, 53, 153, 3815, 423, 46]

```

Train neural network for NLP

```

train_padded = pad_sequences(train_sequences, maxlen=max_length, padding=padding_type, truncating=trunc_type)
print(len(train_sequences[0]))
print(len(train_padded[0]))

print(len(train_sequences[1]))
print(len(train_padded[1]))

print(len(train_sequences[10]))
print(len(train_padded[10]))

16
200
6
200
14
200

```

```
print(train_padded[10])
```

```
[ 8 190 37 201 30 260 293 991 222 53 153 3815 423 46
 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0]
```

```
validation_sequences = tokenizer.texts_to_sequences(validation_articles)
```

```
validation_padded = pad_sequences(validation_sequences, maxlen=max_length, padding=padding_type,
truncating=trunc_type)
```

```
print(len(validation_sequences))
```

```
print(validation_padded.shape)
```

```
1115
```

```
(1115, 200)
```

```
label_tokenizer = Tokenizer()
```

```
label_tokenizer.fit_on_texts(labels)
```

```
training_label_seq = np.array(label_tokenizer.texts_to_sequences(train_labels))
```

```
validation_label_seq = np.array(label_tokenizer.texts_to_sequences(validation_labels))
```

```
print(training_label_seq[0])
```

```
print(training_label_seq[1])
```

```
print(training_label_seq[2])
```

```
print(training_label_seq.shape)
```

```
print(validation_label_seq[0])
```

```
print(validation_label_seq[1])
```

```
print(validation_label_seq[2])
print(validation_label_seq.shape)
```

```
[1]
[1]
[2]
(4457, 1)
[1]
[2]
[1]
(1115, 1)
```

```
reverse_word_index = dict([(value, key) for (key, value) in word_index.items()])
```

```
def decode_article(text):
    return ' '.join([reverse_word_index.get(i, '?') for i in text])
print(decode_article(train_padded[10]))
print('---')
print(train_articles[10])
```

```
i'm gonna home soon want talk stuff anymore tonight k i've cried enough today
????????????????????????????????????????????????????????????
????????????????????????????????????????????????????????????
????????????????????????????????????????????????????????????
????????????????????????????????????????????????????????????
????????????????????????????????????????????????????????????
```

```
---
```

```
I'm gonna home soon want talk stuff anymore tonight, k? I've cried enough today.
```

To implement LSTM

```
model = tf.keras.Sequential([
    tf.keras.layers.Embedding(vocab_size, embedding_dim),
    tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(embedding_dim)),
    tf.keras.layers.Dense(embedding_dim, activation='relu'),
    tf.keras.layers.Dense(6, activation='softmax')
])
model.summary()
```

```
Model: "sequential"
```

Layer (type)	Output Shape	Param #
--------------	--------------	---------

```

=====
embedding (Embedding)      (None, None, 64)      320000

bidirectional (Bidirectional (None, 128)      66048
l)

dense (Dense)              (None, 64)            8256

dense_1 (Dense)            (None, 6)              390

=====
Total params: 394,694
Trainable params: 394,694
Non-trainable params: 0

```

```
print(set(labels))
```

```
{'spam', 'ham'}
```

```
model.compile(loss='sparse_categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
```

```
num_epochs = 10
```

```
history = model.fit(train_padded, training_label_seq, epochs=num_epochs, validation_data=(validation_padded,
validation_label_seq), verbose=2)
```

```
Epoch 1/10
```

```
140/140 - 37s - loss: 0.3177 - accuracy: 0.9251 - val_loss: 0.0387 - val_accuracy: 0.9830 - 37s/epoch - 265ms/step
```

```
Epoch 2/10
```

```
140/140 - 35s - loss: 0.0310 - accuracy: 0.9915 - val_loss: 0.0318 - val_accuracy: 0.9901 - 35s/epoch - 252ms/step
```

```
Epoch 3/10
```

```
140/140 - 32s - loss: 0.0130 - accuracy: 0.9975 - val_loss: 0.0627 - val_accuracy: 0.9857 - 32s/epoch - 230ms/step
```

```
Epoch 4/10
```

```
140/140 - 31s - loss: 0.0060 - accuracy: 0.9987 - val_loss: 0.0478 - val_accuracy: 0.9901 - 31s/epoch - 220ms/step
```

```
Epoch 5/10
```

```
140/140 - 30s - loss: 0.0042 - accuracy: 0.9989 - val_loss: 0.0613 - val_accuracy: 0.9883 - 30s/epoch - 215ms/step
```

```
Epoch 6/10
```

```
140/140 - 29s - loss: 0.0033 - accuracy: 0.9991 - val_loss: 0.0728 - val_accuracy: 0.9883 - 29s/epoch - 210ms/step
```

```
Epoch 7/10
```

```
140/140 - 29s - loss: 0.0020 - accuracy: 0.9996 - val_loss: 0.0540 - val_accuracy: 0.9865 - 29s/epoch - 208ms/step
```

```
Epoch 8/10
```

140/140 - 31s - loss: 7.6466e-04 - accuracy: 0.9998 - val_loss: 0.0644 - val_accuracy: 0.9901 - 31s/epoch - 219ms/step
Epoch 9/10

140/140 - 30s - loss: 3.9159e-04 - accuracy: 1.0000 - val_loss: 0.0678 - val_accuracy: 0.9883 - 30s/epoch - 211ms/step
Epoch 10/10

140/140 - 29s - loss: 1.7514e-04 - accuracy: 1.0000 - val_loss: 0.0726 - val_accuracy: 0.9883 - 29s/epoch - 208ms/step

```
def plot_graphs(history, string):  
    plt.plot(history.history[string])  
    plt.plot(history.history['val_'+string])  
    plt.xlabel("Epochs")  
    plt.ylabel(string)  
    plt.legend([string, 'val_'+string])  
    plt.show()
```

```
plot_graphs(history, "accuracy")
```

```
plot_graphs(history, "loss")
```